

AgentLab — Flowchart Overview

1 ▪ High-Level Agent Roles + Primary Tools

| Role | Core Responsibilities | LLM Prompt “Persona” |
|------------------------------|---|-------------------------------------|
| Student CEO (Human) | Sets vision, approves all outputs | n/a |
| Product Manager Agent | Frame problem, craft user stories, prioritise backlog | “Vision-driven, detail-oriented PM” |
| Design Agent | Wireframes, UX copy, accessibility checks | “Human-centred product designer” |
| Development Agent | Generate code scaffolds, API contracts, tests | “Pragmatic full-stack engineer” |
| AI / Data Agent | Prompt design, ML pipeline, embeddings search | “LLM & data science architect” |
| DevOps Agent | CI/CD, containerisation, | “Automation-first DevOps lead” |

| Role | Core Responsibilities | LLM Prompt “Persona” |
|---------------------------------|---|-----------------------------------|
| | environment configs | |
| Marketing / Growth Agent | Positioning, pitch deck, GTM OKRs | “Story-teller growth hacker” |
| Content Writer Agent | Blog posts, release notes, reflection doc | “Conversational technical writer” |

Tip → Keep prompts modular so you can swap GPT-4o, Claude, or HuggingFace models without rewriting logic.

2 ▪ *Interaction Flow (text version)*

1. **Student enters idea →** routed to **PM Agent**.
2. PM asks clarifying Qs; writes *initial PRD* into memory store.
3. Design Agent consumes PRD → returns Figma wireframe links.
4. Student approves wireframes; PM finalises *MVP user stories*.
5. Development Agent scaffolds repo & OpenAPI spec; pushes to GitHub.
6. AI/Data Agent inserts prompt-wrapper modules & vector DB schema.
7. DevOps Agent spins up preview on Fly.io; URL posted back.
8. Marketing Agent drafts landing-page copy & slide deck.
9. Student reviews; loops back to the relevant agent for revisions.
10. When “ MVP accepted” flag set, Content Writer posts recap article.

Communication is mediated by an **Orchestrator Router** that:

- tags messages with role, thread_id, priority
- enforces max 3-turn cycles before requiring student approval.

3 ▪ Memory Management

| Layer | Contents | Persistence | Retrieval | Extra Notes |
|------------------------|--|---|--|---|
| <i>Ephemeral</i> | Last 10 dialogue turns / agent | In-RAM | Concatenate directly into the next LLM prompt | Cheap & fast; cleared when server restarts. |
| <i>Session Summary</i> | Agent-written bullet digest every ~20 turns | Commit a JSON file to /sessions/YYYY-MM-DD-HH-MM.json on each checkpoint | Use GitHub REST API → fetch the latest file by filename pattern | Optionally trigger commit via a GitHub Action or simple git CLI call. |
| <i>Long-Term</i> | Key-value: decision-log, design-rationale, pivots | 1. Store raw notes in /archive/.../*.json. 2. Nightly GitHub Action generates vector_index.parquet (or .sql) and commits it. | Pull vector_index.*, load into local FAISS / pgvector, run cosine-similarity search, then fetch matching | Keeps repo readable while still enabling semantic search. |

| Layer | Contents | Persistence | Retrieval | Extra Notes |
|-----------------------|--|---|------------------------------------|----------------------------------|
| | | | raw JSON blobs. | |
| <i>Reflection Log</i> | Weekly “what went well / next” self-evaluation | Markdown files under /reflection/week-NN.md | Downloaded raw file via GitHub CDN | Great for peer review & grading. |

Garbage collection.

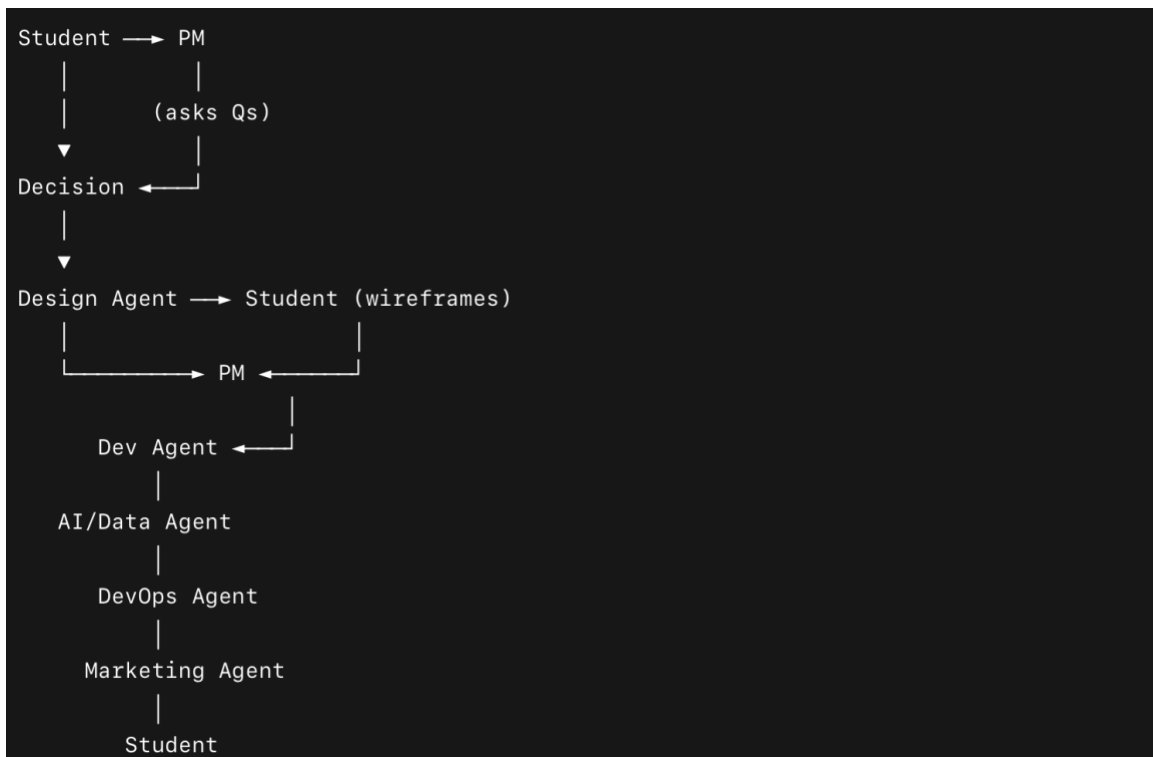
Add a scheduled GitHub Action to delete or archive vectors & JSON older than 90 days to keep the repo slim.

4 ▪ Tools & Integrations

- ❑ **LangChain / CrewAI (or Google ADK / AutoAgent)** – routes messages between agents, manages tool calls, retries, and step ordering.
- ❑ **OpenAI GPT-4o / Claude 3** – core LLMs for reasoning, code generation, and content drafting.
- ❑ **GitHub (private repo)** – long-term storage for JSON session snapshots, Markdown reflections, prompt files, and the generated vector index; gives built-in version control and audit history.
- ❑ **GitHub Actions** – automates commits, nightly embedding jobs that build `vector_index.parquet`, and CI/CD pipelines triggered by the DevOps Agent.
- ❑ **FAISS (or Chroma) local index** – loads the vector file from the repo at app start-up and serves fast cosine-similarity searches without an external vector-DB bill.
- ❑ **Next.js + Tailwind** – responsive chat UI and canvas pane where students view agent messages, wireframes, and code snippets.

- **Figma Plugin API** – streams live wireframe thumbnails or links directly into the chat, keeping design iterations visible.
- **Fly.io or Render** – one-command Docker deploy targets so students can push the whole stack (API + UI) without managing servers.
- **Object Storage (optional: Git LFS, GitHub Releases, or S3)** – holds large binary assets (e.g., high-res design exports) if they exceed normal repo size limits.

5 ▪ *Detailed Interaction Model (state diagram style)*



Edges represent message packets with JSON envelope:

{sender, recipient, content, step_id, status}

6 ■ Example Scenario

| Step | Exchange | Key Output |
|------|---|-------------------------------|
| 1 | Student: “Build campus carpool app.” | Idea registered |
| 2 | PM→Student: “Main pain-point?” | Clarified: parking scarcity |
| 3 | Design Agent → Student: Wireframe v0.1 (Figma) | Visual validated |
| 4 | Dev Agent pushes carpool-app repo | React-Native + Supabase setup |
| 5 | AI/Data Agent adds <i>RideMatchLLM.py</i> | Uses trip vectors + GPT-4o |
| 6 | DevOps Agent deploys to TestFlight | Beta link |
| 7 | Marketing Agent posts deck & tagline: “Park less, ride together.” | GTM asset |
| 8 | Student toggles MVP_ACCEPTED = true | System locks artefacts |

7 ▪ *Key Design Tips (expanded)*

1. **Single-responsibility prompts** → prevents cross-talk.
2. **Conversation IDs** per feature branch to isolate context.
3. **Automated schema validation** on JSON output from agents.
4. **“Human-in-the-loop” gates** after any destructive operation (DB migrations, prod deploy).
5. **Prompt version control** with semantic diff to track performance drift.
6. **Telemetry hooks**—log token usage, latency, agent success/failure rates for course analytics.

8 ▪ *Conclusion*

By combining specialised LLM agents, a clear routing protocol, and a persist-and-retrieve memory layer, *AI Expert TeamMate* turns BADM 350 students into CEOs of a guided startup simulation. The framework’s modular tools (LangChain, Supabase, GitHub Actions) ensure it can scale from a classroom demo to capstone-level, production-ready projects.